



# **SUPERCOMPUTING**

## **@ GEORGIA TECH**

**Attendance**



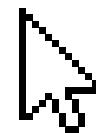
**Club Linktree**





# WORKSHOP 2:

## Matrix Multiplication Optimization



Slides by: Panya B., Min L., Saatvik A.



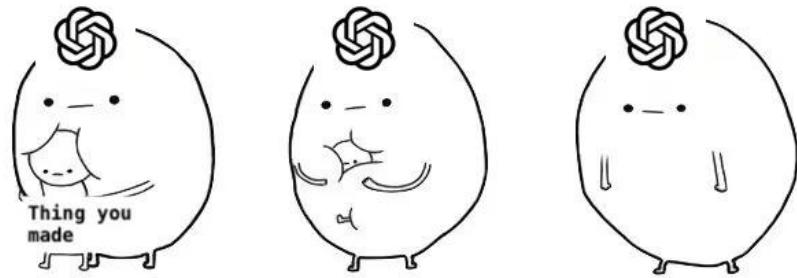
# Agenda

1. Introduction on Matrix Multiplication
2. Tools Overview
3. More on HPC
4. The Message-Passing Paradigm (MPI)
- 5. The Matrix Multiplication Algorithm**
6. Solution 1: Brute Force
7. Optimization 0: Divide-and-conquer
8. Time/Space Complexity Analysis
9. Caching
10. Optimization 1: Tiling
11. Optimization 2: Cannon's Algorithm
12. Optimization 3: Sparse Matrix Multiplication
- 13. Implementation Time!**
14. Applications
15. Key Takeaways



# Why matrix multiplication?

- Mathematical concept that most of you guys already know.
- Most classic and straightforward problem to emphasise the importance of HPC.
- Well-studied and foundational algorithms
- Able to be implemented using a variety of HPC paradigms relatively straightforwardly.
- Broad applications in many areas.





# Tools Overview

- Python 3.11
- mpi4py
- PACE Instructional Cluster (ICE)
- Your brains!



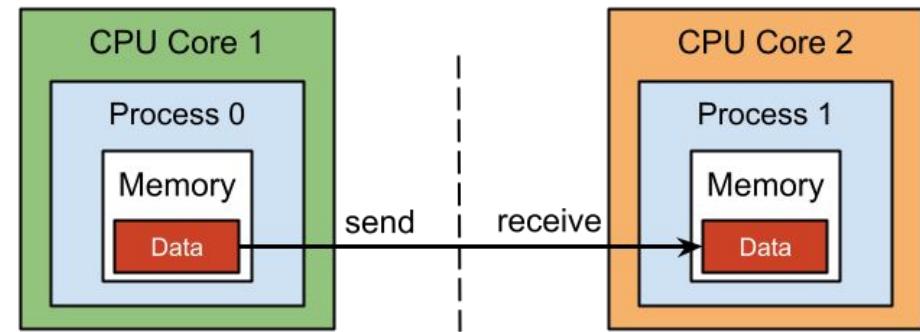
# A little more on HPC...

- Despite what NVIDIA may want you to think, CUDA is not the only software stack that exists in the HPC world!
- There are many programming paradigms in the HPC world, with pros and cons for each.
  - Data-Parallel,
  - Shared-Memory,
  - Message-Passing,
  - ...
- In this workshop we will be focused on message-passing using the MPI protocol!
- MPI allows communication between different machines with their own distributed memory.



# The Message-Passing Paradigm

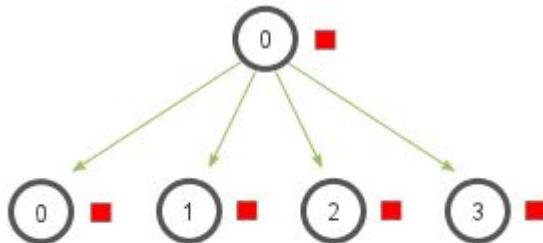
- Group of nodes = “world”
- Each node has a “rank”
- Send and receive messages between different ranks
- Common functions:
  - Broadcast
  - Scatter
  - Reduce / AllReduce
  - Gather / AllGather
  - All-to-all / Many-to-many



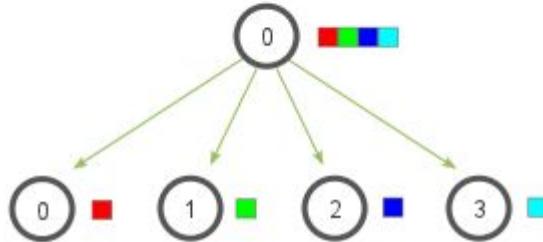


# MPI Functions

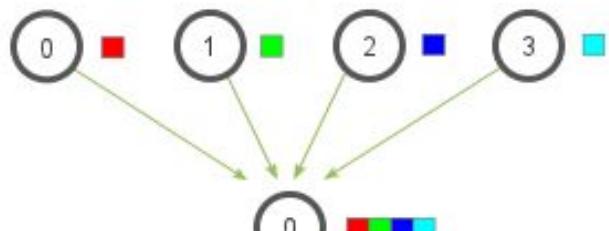
MPI\_Bcast



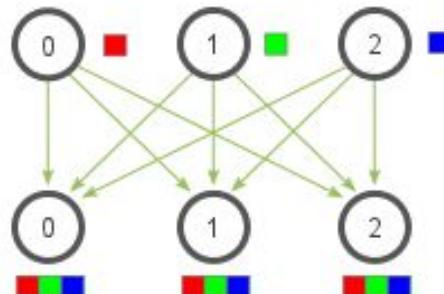
MPI\_Scatter



MPI\_Gather



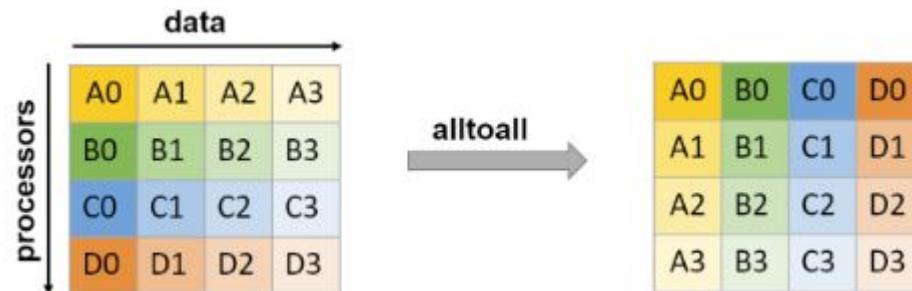
MPI\_Allgather





# MPI Functions

- All-to-All:



- Many-to-many is similar, but with variable data lengths



# mpi4py

**Rank, Size, COMM:** `comm.Get_rank()`, `comm.Get_size()`, `MPI.COMM_WORLD`

**Send:** `comm.send(sendbuf, dest, tag)`

**Receive:** `comm.recv(source, tag)`

**Broadcast:** `comm.bcast(sendbuf, root)`

**Reduce:** `comm.reduce(int, root, operation)`

**AllReduce:** `comm.allreduce(int, operation)`

**Scatter:** `comm.scatter(sendbuf, recvbuf, root=root)`

**Gather:** `comm.gather(sendbuf, recvbuf, root=root)`

**AllGather:** `comm.allgather(sendbuf, recvbuf)`



# Agenda

1. Introduction on Matrix Multiplication
2. Tools Overview
3. More on HPC
4. The Message-Passing Paradigm (MPI)
- 5. The Matrix Multiplication Algorithm**
6. Solution 1: Brute Force
7. Time/Space Complexity Analysis
8. Optimization 0: Divide-and-conquer
9. Caching
10. Optimization 1: Tiling
11. Optimization 2: Cannon's Algorithm
12. Optimization 3: Sparse Matrix Multiplication
- 13. Implementation Time!**
14. Applications
15. Key Takeaways



# Inner Matrix Multiplication Algorithm

**Problem Statement:** Given matrices A and B of size  $(m \times n)$  and  $(n \times p)$  respectively, where,  $m \neq n \neq p$  and  $m, n, p > 0$ , calculate their inner product AB.

$$\text{row } i \begin{bmatrix} \quad \\ \quad \\ \quad \\ \boxed{\quad} \\ \quad \end{bmatrix} \times \begin{bmatrix} \quad \\ \quad \\ \quad \\ \boxed{\quad} \\ \quad \end{bmatrix} = \begin{bmatrix} \quad \\ \quad \\ \quad \\ \boxed{\quad} \\ \quad \end{bmatrix}$$

col  $j$

$$c_{ij} = \sum_{k=1}^n a_{ik} \cdot b_{kj}$$
$$A_{m \times n} \times B_{n \times p} = C_{m \times p}$$



# Solution 1: Brute Force Approach

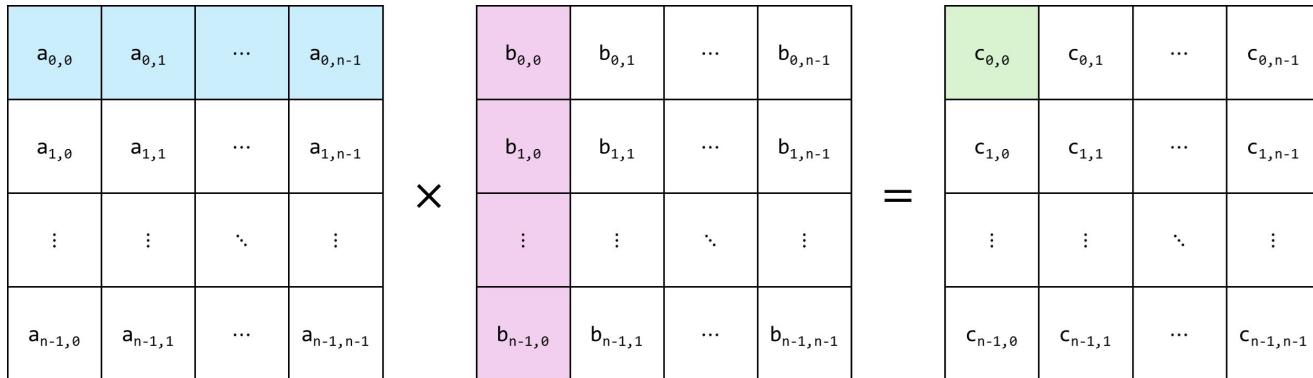
Naive Approach: Using Nested Loops –  $O(n^3)$  Time and  $O(n^2)$  Space

RECTANGULAR-MATRIX-MULTIPLY( $A, B, C, p, q, r$ )

```
1  for  $i = 1$  to  $p$ 
2      for  $j = 1$  to  $r$ 
3          for  $k = 1$  to  $q$ 
4               $c_{ij} = c_{ij} + a_{ik} \cdot b_{kj}$ 
```



# Solution 1: Brute Force Approach



The diagram shows the computation of element  $c_{i,j}$  as a sum of products. It is represented as  $c_{i,j} = \sum a_{i,k} \cdot b_{k,j}$ , where the terms are shown as  $a_{i,0}, a_{i,1}, \dots, a_{i,n-1}$  and  $b_{0,j}, b_{1,j}, \dots, b_{n-1,j}$ .

$a_{i,0}$	$a_{i,1}$	$\dots$	$a_{i,n-1}$
-----------	-----------	---------	-------------

$c_{i,j}$

$=$

$\times$     $\times$     $\times$     $\times$

$b_{0,j}$	$b_{1,j}$	$\dots$	$b_{n-1,j}$
-----------	-----------	---------	-------------

RECTANGULAR-MATRIX-MULTIPLY( $A, B, C, p, q, r$ )

```
1  for  $i = 1$  to  $p$ 
2      for  $j = 1$  to  $r$ 
3          for  $k = 1$  to  $q$ 
4               $c_{ij} = c_{ij} + a_{ik} \cdot b_{kj}$ 
```



# Why do we need parallelization?

- Suppose matrices are of size  $(3 \times 3)$   $\Rightarrow$  total number of instructions = 27
- Suppose matrices are of size  $(10^6 \times 10^6)$   $\Rightarrow$  total number of instructions =  $10^{18}$



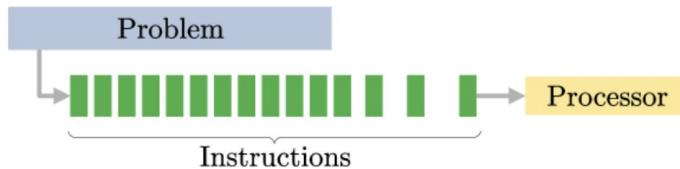
How do we make  
this faster?



# Time Complexity Analysis of Parallel Programs

## Serial Programs

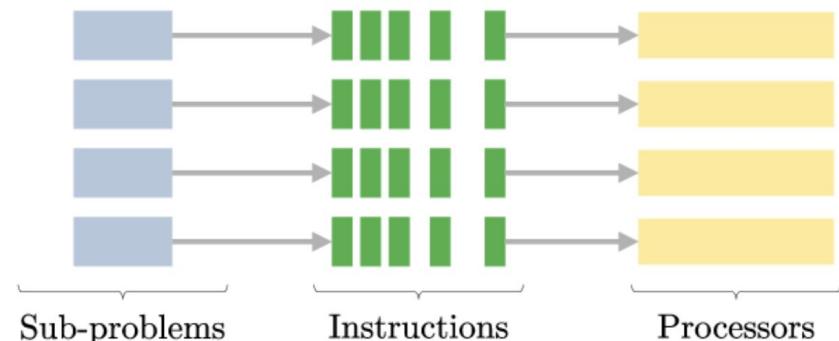
amount of work done by the program  $=$  time it takes for the program to finish.



## Parallel Programs

amount of work done by the program  $\neq$  time it takes for the program to finish!

VS





# Optimization 0: Divide and Conquer

- Theorem: Suppose that we partition each of  $A$ ,  $B$ , and  $C$  into four  $n/2 \times n/2$  matrices

$$A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}, \quad B = \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}, \quad C = \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix}, \quad (4.9)$$

so that we rewrite the equation  $C = A \cdot B$  as

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \cdot \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}. \quad (4.10)$$

Equation (4.10) corresponds to the four equations

$$C_{11} = A_{11} \cdot B_{11} + A_{12} \cdot B_{21}, \quad (4.11)$$

$$C_{12} = A_{11} \cdot B_{12} + A_{12} \cdot B_{22}, \quad (4.12)$$

$$C_{21} = A_{21} \cdot B_{11} + A_{22} \cdot B_{21}, \quad (4.13)$$

$$C_{22} = A_{21} \cdot B_{12} + A_{22} \cdot B_{22}. \quad (4.14)$$

Cormen et al.

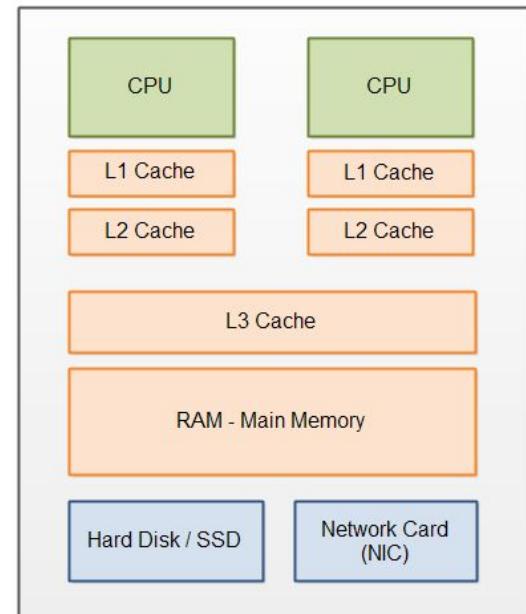
- This can be done on a single processor using the divide-and-conquer paradigm, or alternatively, we can split the workload across multiple processors!



# Caching

- Memory access is expensive
- **Caches** are fast, smaller memory near the CPU that stores recently or frequently accessed data
- **Spatial locality:** processor likely to access memory locations near each other
- **Temporal locality:** memory accessed at a specific time is likely to be accessed again
- **Problem with naive matmul:** poor spatial and temporal locality → lots of cache misses

What if we broke the matrices into smaller blocks stored in the cache ...





# Optimization 1: Blocked Multiplication

**Goal:** Increase parallelism

**How?**

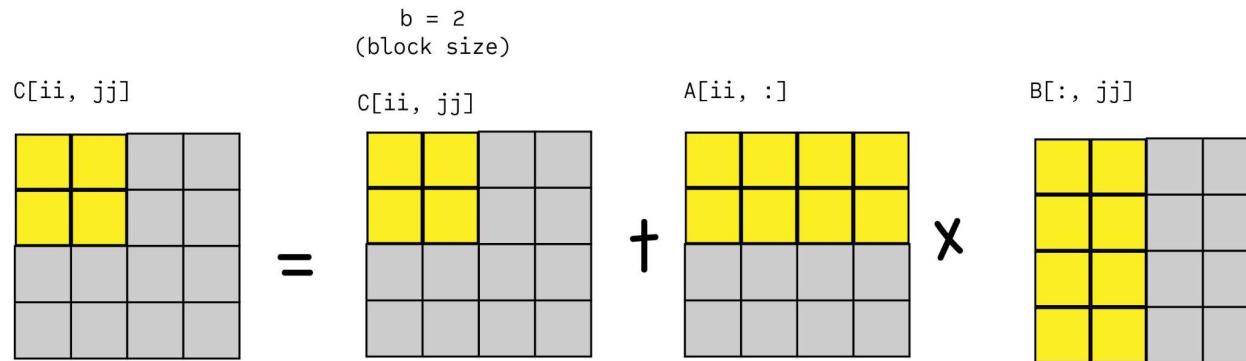
- Divide the matrix into evenly sized blocks, 1 per processor (2D partitioning of both matrices).
- Send all the blocks in a row to all processors in a row, and send all the blocks in a column to all processors in a column.
- Each processor performs matrix multiplication for its block.
- The root collects all of the results.

**Pros and Cons?**

- Computation is divided among the processors.
- However, total cache usage is still suboptimal as blocks are present multiple times.



# Optimization 1: Blocked Multiplication



ii, jj, kk denote block indices while i, j, k denote element indices

```
// C = C + A * B
// b = n / N (where b is the block size)
for ii = 1 to N:
    for jj = 1 to N:
        for kk = 1 to N:
            C[ii, jj] += A[ii, kk] * B[kk, jj]
```



# Optimization 2: Cannon's Algorithm

$$c_{ij} = \sum_{k=1}^n a_{ik} \cdot b_{kj}$$

## Key Observation:

- Partition A and B in  $p$  square blocks, denoted  $A_{ij}$  and  $B_{ij}$ .
- Computing  $C_{ij}$  just needs all submatrices  $A_{ik}$  and  $B_{kj}$  for  $0 \leq k \leq \sqrt{p}$ .

$$C_{i,j} = \sum_{k=0}^{\sqrt{p}-1} A_{i,(i+j+k)\% \sqrt{p}} B_{(i+j+k)\% \sqrt{p},j}$$



# Optimization 2: Cannon's Algorithm

A(0,0)	A(0,1)	A(0,2)
A(1,0)	A(1,1)	A(1,2)
A(2,0)	A(2,1)	A(2,2)

A(0,0)	A(0,1)	A(0,2)
A(1,1)	A(1,2)	A(1,0)
A(2,2)	A(2,0)	A(2,1)

A(0,1)	A(0,2)	A(0,0)
A(1,2)	A(1,0)	A(1,1)
A(2,0)	A(2,1)	A(2,2)

A(0,2)	A(0,0)	A(0,1)
A(1,0)	A(1,1)	A(1,2)
A(2,1)	A(2,2)	A(2,0)

B(0,0)	B(0,1)	B(0,2)
B(1,0)	B(1,1)	B(1,2)
B(2,0)	B(2,1)	B(2,2)

B(0,0)	B(1,1)	B(2,2)
B(1,0)	B(2,1)	B(0,2)
B(2,0)	B(0,1)	B(1,2)

B(1,0)	B(2,1)	B(0,2)
B(2,0)	B(0,1)	B(1,2)
B(0,0)	B(1,1)	B(2,2)

B(2,0)	B(0,1)	B(1,2)
B(0,0)	B(1,1)	B(2,2)
B(1,0)	B(2,1)	B(0,2)

Initial A, B

A, B initial  
alignment

A, B after  
shift step 1

A, B after  
shift step 2



# Optimization 2: Cannon's Algorithm

**Goal:** Improve memory efficiency

**How?**

- Multiplies 2 square matrices using processors in a 2D grid.
- Distributes submatrices (blocks) of A and B to each processor.
- Minimizes communication by reusing data blocks over multiple computation steps.

**Key Steps**

- **Align data:** Shift matrix A's blocks left by row index i. Shift matrix B's blocks up by column index j.
- **Main Loop** (repeat  $\sqrt{p}$  times): Each processor multiplies local  $A_{ij} \times B_{ij}$  and accumulates the result in  $C_{ij}$ . Then:
  - Shift all blocks of A left by 1 (wraparound)
  - Shift all blocks of B up by 1 (wraparound)



# Optimization 2: Cannon's Algorithm

**What improvements does this have?**

- Avoids global all-to-all broadcast (reduces computation time)
- Computation dominates for larger matrices
- Memory needed does not increase as processors increase



# Optimization 3: Sparse Matrix Mult.

- **Sparse Matrix:** Matrix with primarily zeros.
- Much of matrix multiplication done in HPC is SpMM.
- Sometimes using dense matrix multiplication algorithms is not necessary and wastes time and memory.
- Notice that (nonzero x nonzero) calculations will occur much lesser in frequency than multiplications that involve 0s.
- ⇒ Optimize matmul by making sure we only do (nonzero x nonzero) multiplications!
- Use a better data structure to ensure this!

$$c_i \sum_{a_{iv} \neq 0} a_{iv} b_v. \text{ for } 1 \leq i \leq p.$$

“The  $i$ th row of resulting matrix  $C$  is the linear combination of the  $v$  rows of  $B$  for which  $a_{iv} \neq 0$ .  
– Gustavson (September, 1978)



# Designing A SpMM Algorithm

Step 1. Designing a better data structure:

→ One Solution: The **Coordinate Representation (COO)**

	0	1	2	3
0	12		26	
1				
2		19		
3		14		7

(0, 0, 12)  
(0, 2, 26)  
(2, 1, 19)  
(3, 1, 14)  
(3, 3, 7)

row + column index +  
weight per nonzero  
(easy to build / modify)



# Designing A SpMM Algorithm

**Step 2.** Minimize the number of multiplications we are doing with 0 elements by using the dot product formula.

$$c_i \sum_{a_{iv} \neq 0} a_{iv} b_v. \quad \text{for } 1 \leq i \leq p.$$

"The  $i$ th row of resulting matrix  $C$  is the linear combination of the  $v$  rows of  $B$  for which  $a_{iv} \neq 0$ ."

- Gustavson (September, 1978)

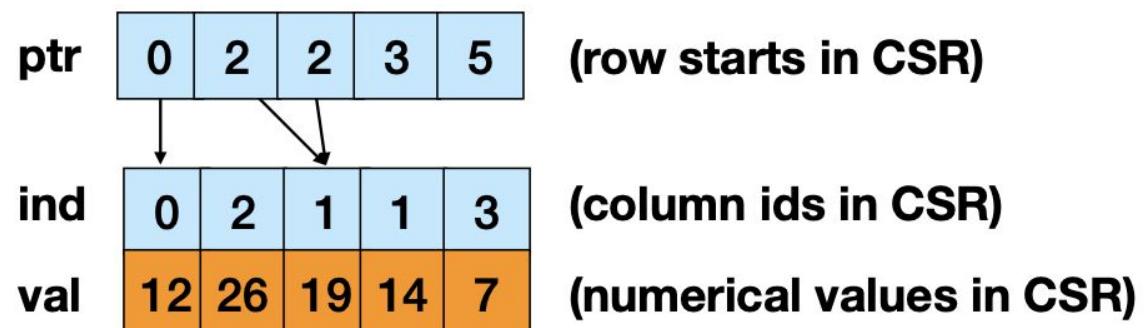
```
RowWISE_SPGEMM(C, A, B)
1 // set matrix C to Ø
2 for  $a_{i*}$  in matrix A in parallel
3   do for  $a_{ik}$  in row  $a_{i*}$ 
4     do for  $b_{kj}$  in row  $b_{k*}$ 
5       do value  $\leftarrow a_{ik} b_{kj}$ 
6       if  $c_{ij} \notin c_{i*}$ 
7         then insert ( $c_{ij} \leftarrow value$ )
8       else  $c_{ij} \leftarrow c_{ij} + value$ 
```



# But wait... Limitations of COO

- Revisit step 1: Notice that the only way to implement COO data structure results in **low cache locality!**
- Solution: A new data structure called **Compressed Sparse Row (CSR) representation.**

	0	1	2	3
0	12		26	
1				
2		19		
3		14		7





# Implementation Time!



# Set Up

Getting Started with PACE



# SSH Instructions

- **IMPORTANT:** Make sure that you are connected to the GT VPN!
- Connect to Host: For everyone, the host name is  
`<gt-username>@login-ice.pace.gatech.edu`
- Type in your GT password upon seeing the prompt. DO NOT BE ALARMED IF NOTHING SHOWS UP AS YOU TYPE (its for safety etc.) Just finish typing your password and press <Enter>.

```
(base) ipsec-10-2-128-166:~ panyabhinder$ ssh pbhinder3@login-ice.pace.gatech.edu
pbhinder3@login-ice.pace.gatech.edu's password: 
```

[https://gatech.service-now.com/home?id=kb\\_article\\_view&sysparm\\_article=KB0042100](https://gatech.service-now.com/home?id=kb_article_view&sysparm_article=KB0042100)



# Remote SSH w/ Terminal

- Please check that you have access to an **SSH-Capable Terminal!**
  - Windows Powershell
  - MacOS Terminal
  - Linux Terminal
- We'll use **SSH** to log onto the cluster you'll be using today.
- Follow the previous instructions to login to the PACE cluster. If you've successfully logged in, you'll see a prompt like the following:

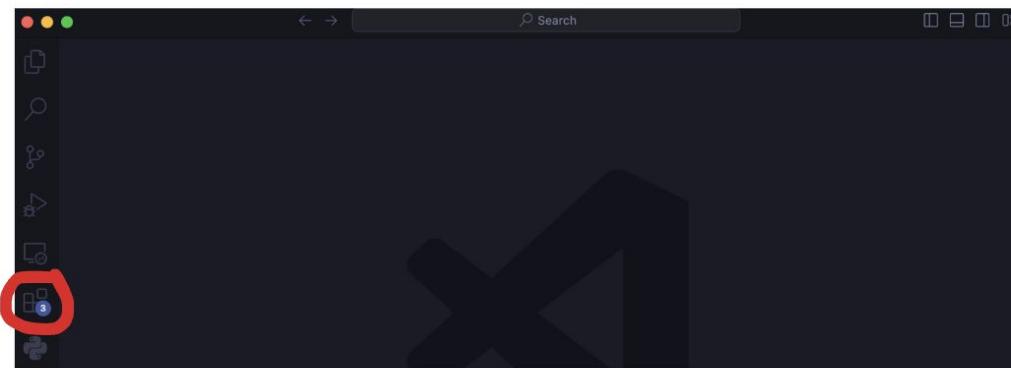
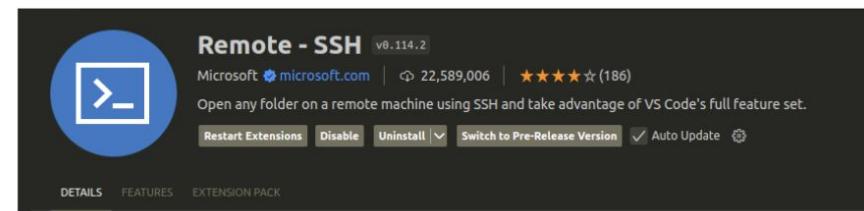
```
[pbhinder3@login-ice-3 ~]$ █
```

- You are now in a **login node** of the cluster, which you can use to access files, write code, submit jobs, etc.



# Remote SSH w/ VSCode

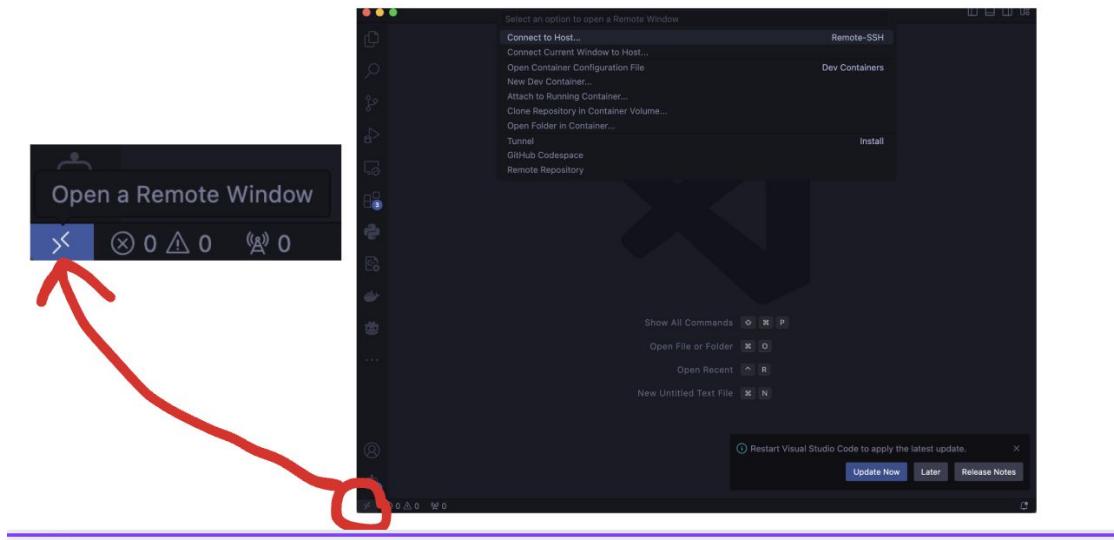
Install the extension from the VSCode extension store.





# Remote SSH w/ VSCode

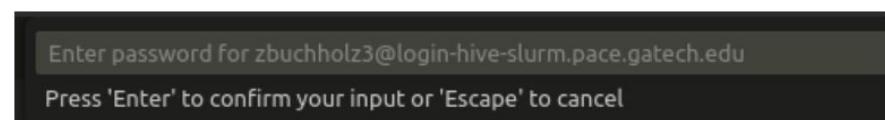
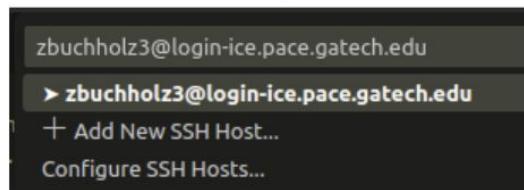
2. Notice that at the bottom left of the application, you will now see a button which appears to show a broken link. Click on it.



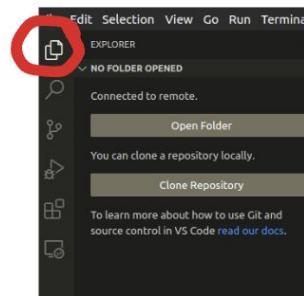


# Remote SSH w/ VSCode

3. Select Connect-to-Host (should be the first option) and enter the hostname and password as shown previously.



4. You should now be able to open the folders icon to see the contents of the remote.





# Set up repository

In your terminal, clone the workshop repository:

```
git clone https://github.com/suco-gt/Parallel-Optimization-Workshop.git
```

Follow the setup instructions in the README:

- Run `chmod +x *.sh` to give executable permissions to the shell scripts that are going to be used for this workshop.
- Run `./get_nodes.sh [num_nodes]` to allocate nodes on the cluster for 1 hour. Default num\_nodes is 4.
- `module load anaconda3`
- Install conda environment: `conda create --name sp25_suco2 -c conda-forge python=3.11 numpy openmpi mpi4py -y`
- `conda activate sp25_suco2`
- Test if mpi4py is working for you successfully by running `./run.sh test`.
- You will run your code with the given bash file “run.sh”:
- For example: To run the intro exercises- “`./run.sh intro`”



# Your Job

**Your tasks (choose at least 1):**

Complete the introduction to MPI problems in `0-mpi-intro.py`

**OR**

Race to implement each optimization! (Feel free to work in groups)

1. Implement the brute force matmul implementation in `bruteforce.py`
2. Implement the blocked matrix multiplication optimization in `blocked.py`
3. Implement Cannon's Algorithm in `cannon.py`
4. Implement COO SpGEMM in `spgemm.py`

Let us know if you have any questions!

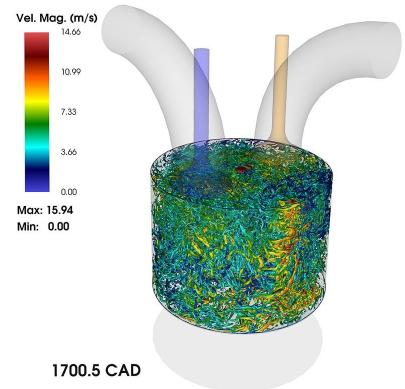
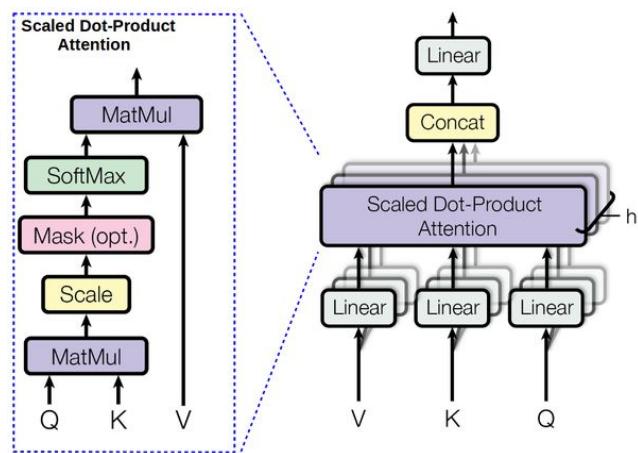
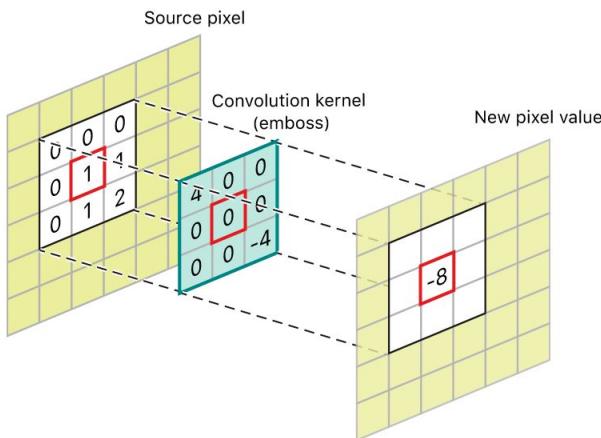


# Agenda

1. Introduction on Matrix Multiplication
2. Tools Overview
3. More on HPC
4. The Message-Passing Paradigm (MPI)
- 5. The Matrix Multiplication Algorithm**
6. Solution 1: Brute Force
7. Time/Space Complexity Analysis
8. Optimization 0: Divide-and-conquer
9. Caching
10. Optimization 1: Tiling
11. Optimization 2: Cannon's Algorithm
12. Optimization 3: Sparse Matrix Multiplication
- 13. Implementation Time!**
- 14. Applications**
15. Key Takeaways



# Applications





# Key Takeaways

- Parallelism and scaling
  - Breaking up computation into smaller tasks speed up execution
  - Difference between shared memory (OpenMP) and distributed memory parallelism (MPI)
- Memory hierarchy and data locality - hardware impacts efficiency
  - Tiling and blocking improves cache efficiency
  - Impact of sparsity on performance
- Algorithmic trade-off
  - Communication vs computation bottlenecks
  - Choosing right optimization based on problem size and hardware constraints